

# Probabilistic Graphical Models

## Lecture 13: Learning Graphical Models and Parameter Estimation

Matthew Walter

TTI-Chicago

May 26, 2020

- Graphical models as **representations** of probability distributions
  - Directed graphical models (i.e., Bayesian networks)
  - Undirected graphical models (i.e., Markov random fields, conditional random fields, and factor graphs)
- Algorithms for performing **inference** in these networks
  - Variable elimination (sum-product)
  - Exact inference in clique trees via message passing
  - Exact MAP inference (max-product)
  - Approximate inference via variational methods
  - Approximate inference via sampling
- Techniques for **learning** graphical model parameters and structure
  - Learning Bayesian networks
  - Learning Markov random fields
  - Learning with incomplete data

# How Can We Acquire a Model?

- We have several options:
  - Rely on expert knowledge to determine graph structure and parameters
  - Use learning to determine the potentials, i.e., **parameter learning**
  - Use learning to determine the graph structure, i.e., **structure learning**
- Manual design is difficult and can require a lot of time
- Often, we have access to a set of examples (samples) from the distribution that we want to model

# How Can We Acquire a Model?

- Assume that the domain is governed by an underlying distribution  $P^*$  that is induced by some network model  $\mathcal{M}^* = (G^*, \theta^*)$
- We are given a dataset  $\mathcal{D}$  of  $M$  samples from  $P^*$
- We assume that the samples are **independent and identically distributed (IID)**
- We have a family of models  $\{\mathcal{M}\}$  (given or specified by us), and our task is to learn some model  $\hat{\mathcal{M}} \in \{\mathcal{M}\}$  that defines a distribution  $P_{\hat{\mathcal{M}}}$
- We can learn model parameters for a fixed structure, or both the structure and parameters

# Goals of Learning

- The goal of learning is to return a model  $\hat{\mathcal{M}}$  that precisely captures the distribution  $P^*$  from which our data was sampled
- In general, this is not achievable due to
  - Access to a small amount of data relative to the number of random variables, providing a sparse sampling of the true distribution
  - Computational reasons
- Consequently, we need to select  $\hat{\mathcal{M}}$  that constructs the “best” approximation to  $\mathcal{M}^*$
- Raises the question: How do we define “best”?

# How Do We Define “Best”?

- The measure of “best” depends on the specific goals
  - ① Density estimation: We want an estimate of the distribution  $\hat{P}$  that is as close as possible to  $P^*$  (e.g., for general inference)
  - ② Specific prediction task: We will use the model to make a particular prediction (e.g., classification, segmentation, or depth estimation)
  - ③ Knowledge discovery: We are interested in the structure of the model

- We want to learn the full distribution, so that we can answer *any* probabilistic query
- The learning problem as one of **density estimation**
- We want to construct  $\hat{\mathcal{M}}$  such that  $\hat{P}$  is as “close” as possible to  $P^*$
- How do we evaluate “closeness”?
- **KL-divergence** (specifically, the M-projection) is one possibility

$$D(P^* \parallel \hat{P}) = \mathbb{E}_{P^*} \left[ \log \frac{P^*(\mathcal{X})}{\hat{P}(\mathcal{X})} \right]$$

(recall that it measures the loss when representing  $P^*$  with  $\hat{P}$ )

# Density Estimation: Expected Log-Likelihood

- We can simplify the expression as

$$D(P^* \parallel \hat{P}) = \mathbb{E}_{P^*} \left[ \log \frac{P^*(\mathcal{X})}{\hat{P}(\mathcal{X})} \right] = -H_{P^*}(\mathcal{X}) - \mathbb{E}_{P^*} [\log \hat{P}(\mathcal{X})]$$

- The first term on the RHS does not depend on  $\hat{P}$  (and is often unknown since we don't know  $P^*$ )
- Thus, finding the *minimal* M-projection is equivalent to *maximizing* the **expected log-likelihood**

$$\mathbb{E}_{P^*} [\log \hat{P}(\mathcal{X})]$$

- Encourages  $\hat{P}$  to assign high probability to instances likely under  $P^*$
- Because of log scale,  $\hat{P}$  will be reasonably large where  $P^*$  is not small
- Without  $H_{P^*}(\mathcal{X})$ , we can compare different models, but we don't know how close they are from the optimum



# Density Estimation: Maximum Likelihood

- When  $P^*$  is unknown, we approximate the expected log-likelihood

$$\mathbb{E}_{P^*}[\log \hat{P}(\mathcal{X})]$$

with the *empirical* log-likelihood

$$\mathbb{E}_{\mathcal{D}}[\log \hat{P}(\mathcal{X})] = \frac{1}{|\mathcal{D}|} \sum_{\boldsymbol{\xi} \in \mathcal{D}} \log \hat{P}(\boldsymbol{\xi})$$

for some set of samples  $\mathcal{D} = \{\boldsymbol{\xi}^{(1)}, \boldsymbol{\xi}^{(2)}, \dots, \boldsymbol{\xi}^{(M)}\}$

- Gives rise to maximum likelihood learning:

$$\max_{\hat{\mathcal{M}}} \frac{1}{|\mathcal{D}|} \sum_{\boldsymbol{\xi} \in \mathcal{D}} \log \hat{P}(\boldsymbol{\xi})$$

# Likelihood, Loss, and Risk

- A **loss function**  $\text{loss}(\boldsymbol{x}, \mathcal{M})$  measures the loss that a model  $\mathcal{M}$  makes on a particular instance  $\boldsymbol{x}$
- When instances are sampled from  $P^*$ , our objective is to minimize the **expected loss** (aka the **expected risk**)

$$\mathbb{E}_{P^*}[\text{loss}(\boldsymbol{x}, \mathcal{M})]$$

- In the case of density estimation, we are interested in log-loss

$$\text{loss}(\boldsymbol{x}, \hat{\mathcal{M}}) = -\log \hat{P}(\boldsymbol{x})$$

- Since  $P^*$  is unknown, we approximate the expectation using the empirical average, i.e., the **empirical risk**

$$\mathbb{E}_{\mathcal{D}}[\text{loss}(\mathcal{X}, \mathcal{M})] = \frac{1}{|\mathcal{D}|} \sum_{\boldsymbol{\xi} \in \mathcal{D}} \text{loss}(\boldsymbol{\xi}, \hat{\mathcal{M}})$$

# Prediction Tasks

- Density estimation is useful for performing probabilistic queries over an arbitrary subset of the random variables
- However, we are often interested in performing prediction over a subset of variables  $\mathbf{Y}$  given some others  $\mathbf{X}$  (e.g., classification tasks, such as image segmentation or document classification)
- Rather than learn the entire joint distribution, we can focus specifically on the conditional distribution  $P(\mathbf{Y} | \mathbf{X})$  (recall CRFs)
- We can then consider several different loss functions
  - Classification error (aka 0/1 loss)
  - Hamming loss (fraction of wrong predictions)
  - Conditional log-loss:  $\mathbb{E}_{P^*} \left[ \log \hat{P}(\mathbf{Y} | \mathbf{X}) \right]$  (analogous to log-loss)

# Prediction Tasks: Conditional Log-Likelihood

- We concentrate on predicting the conditional distribution  $P(\mathbf{Y} | \mathbf{X})$
- Conditional log-loss measures how well the learned model predicts  $\mathbf{Y}$  given  $\mathbf{X}$  in the data, but not the distribution over  $\mathbf{X}$

$$\text{loss}(\mathbf{x}, \mathbf{y}, \hat{\mathcal{M}}) = -\log \hat{P}(\mathbf{y} | \mathbf{x})$$

- This is the objective used to train conditional random fields (CRFs), such as the ones that we saw earlier

**input:** two images



**output:** disparity



# Prediction Tasks: Structured Prediction

- In **structured prediction**, we predict  $\mathbf{y}$  given  $\mathbf{x}$  as

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y}} P(\mathbf{y} | \mathbf{x})$$

- One reasonable choice for the loss function is the **classification error**

$$\mathbb{E}_{(\mathbf{x}, \mathbf{y}) \sim P^*} \left[ \mathbb{1}(\exists \mathbf{y}' \neq \mathbf{y} \text{ s.t. } \hat{P}(\mathbf{y}' | \mathbf{x}) \geq \hat{P}(\mathbf{y} | \mathbf{x})) \right]$$

(i.e., probability over all  $(\mathbf{x}, \mathbf{y})$  that we predict the wrong assignment)

# Empirical Risk and Overfitting

- Given data  $\mathcal{D}$  sampled from  $P^*$ , define the *empirical distribution* as

$$\bar{P}_{\mathcal{D}}(A) = \frac{1}{M} \sum_m \mathbb{1}(\xi^{(m)} \in A)$$

- If  $\mathcal{D} = \{\xi^{(1)}, \xi^{(2)}, \dots, \xi^{(M)}\}$  is a sequence of IID samples  $\xi^{(m)} \sim P^*$ ,

$$\lim_{M \rightarrow \infty} \bar{P}_{\mathcal{D}}(A) = P^*(A)$$

- $\bar{P}_{\mathcal{D}}$  is the distribution that maximizes the log-likelihood of the data
- Optimizing empirical risk can be sensitive to **overfitting** to the data
- Consider a domain with 100 binary variables ( $2^{100}$  assignments):
  - Suppose that  $\mathcal{D}$  contains 1000 instances (likely distinct)
  - We will estimate  $\hat{P}$  that assigns 0.001 to each of the 1000 assignments, and 0 to the  $2^{100} - 1000$  others
- Alternatively, consider a Bayesian network with some random variables having a large number  $k$  of parents
  - Number of parameters for each CPD is exponential in  $k$
  - It is unlikely that  $\mathcal{D}$  will span all instantiations of parents
- In order to improve generalization, we restrict the **hypothesis space**

# Bias-Variance Trade-Off

- If the hypothesis space is very limited, we may not be able to learn  $P^*$ , even with unlimited data
- This introduces a **bias** in how close learning can approximate the true distribution
- If we select a highly expressive hypothesis space, we can better represent the data
- When we have a small amount of data, multiple models may fit well, possibly even better than the true model
- Small perturbations of the data will result in very different estimates, i.e., high **variance**
- There is an inherent **bias-variance trade-off** when selecting hypothesis class
- Error in learning due to both bias and variance

# How Can We Avoid Overfitting?

- Impose hard constraints, e.g., limiting hypothesis class
  - Bayesian networks with at most  $d$  parents per node
  - Pairwise MRFs (vs. higher-order potentials)
- Update objective to include soft preference for simpler models via **regularization**

$$\text{objective}(\mathbf{x}, \mathcal{M}) = \text{loss}(\mathbf{x}, \mathcal{M}) + R(\mathcal{M})$$

- Can evaluate generalization performance via cross-validation



# Learning Procedure

- We assume input of the form:
  - ① Prior knowledge and/or constraints on the model class  $\hat{\mathcal{M}}$
  - ② A set  $\mathcal{D} = \{\xi^{(1)}, \xi^{(2)}, \dots, \xi^{(M)}\}$  of IID samples from  $P^*$
- The output is a model  $\hat{\mathcal{M}}$  that may include the structure and/or parameters of the graphical model
- The specifics of a particular learning algorithm vary with
  - ① The type of output, i.e., a Bayesian network or Markov random field
  - ② The constraints that we place on  $\hat{\mathcal{M}}$
  - ③ The extent to which the training data is fully observed

# Model Constraints

- Constraints on the model define the hypothesis space that specifies the class of admissible models
- These constraints may be over the structure, the parameters, or both
  - ① We may be given the graph structure and tasked with learning some or all of the parameters
  - ② We may not know the structure or parameters, and need to learn both
  - ③ We may not even know the complete set of variables over which the distribution  $P^*$  is defined
- As we discussed, there is a trade-off between over- and under-constraining the model space
  - If the hypothesis space is too restricted, it may not contain  $P^*$
  - On the other hand, if the hypothesis space is too large, we may overfit, learning a model that assigns high likelihood to the data, but is a poor approximation to  $P^*$

# Data Observability

- Different learning problems exhibit different amounts of observability with respect to the data:
  - ① Each datapoint  $\xi^{(i)}$  provides an instantiation to every random variable
  - ② Each datapoint  $\xi^{(i)}$  provides an instantiation for a subset of the random variables, though every variable is observed at least once
  - ③ There are some *hidden* variables that are never observed in the training data (i.e., we don't know the full set of random variables)
- Partial observability is encountered in many practical problems (observing all random variables as part of each sample may be difficult or impossible)
- Hidden variables are often useful
  - They help us better understand the nature of the problem (i.e., knowledge discovery)
  - They simplify the distribution and graphical model (e.g., consider a naive Bayes model)

# Summary of Learning<sup>1</sup>

- 1 Decide on an objective and corresponding loss

$$\mathbb{E}_{P^*}[\text{loss}(\mathbf{x}, \mathcal{M})]$$

- 2 Determine how to best estimate this from what we have, e.g., regularized empirical loss

$$\mathbb{E}_{\mathcal{D}}[\text{loss}(\mathbf{x}, \mathcal{M})] + R(\mathcal{M})$$

When used with log-loss, the regularization term can be interpreted as a prior distribution over models,  $P(\mathcal{M}) \propto \exp(-R(\mathcal{M}))$  (called *maximum a posteriori (MAP)* estimation)

- 3 Determine how to optimize over this objective function

$$\min_{\mathcal{M}} \mathbb{E}_{\mathcal{D}}[\text{loss}(\mathbf{x}, \mathcal{M})] + R(\mathcal{M})$$

(We will start by assuming **complete (fully observable)** data)

---

<sup>1</sup>Obviously, learning deserves an entire course (Greg's ML class is excellent.)

# Maximum Likelihood Parameter Estimation

- A function  $P(\mathbf{X}; \boldsymbol{\theta})$  is a *parametric model* (parametric family) specified in terms of a set of parameters  $\boldsymbol{\theta} \in \Theta$  (parameter space)
- If  $x$  is a multinomial taking on  $K$  values, the parameters  $\boldsymbol{\theta} \in \mathbb{R}^K$  define the distribution as:

$$P(x = x^k; \boldsymbol{\theta}) = \theta_k \quad \text{where} \quad \Theta = \left\{ \boldsymbol{\theta} \in [0, 1]^K : \sum_i \theta_i = 1 \right\}$$

# Maximum Likelihood Parameter Estimation

- The *likelihood function* is the probability of the data

$$L(\boldsymbol{\theta} : \mathcal{D}) = \prod_m P(\mathbf{x}^{(m)}; \boldsymbol{\theta})$$

- Consider  $M$  IID tosses of a possibly biased coin  
 $\mathcal{D} = \{x^{(1)}, x^{(2)}, \dots, x^{(M)}\}$ , where  $x^{(i)} \in \{H, T\}$
- The likelihood of the data as a function of the parameter vector  
 $\boldsymbol{\theta} = [\theta, \theta - 1]$ , where  $\theta = P(H)$

$$L(\boldsymbol{\theta} : \mathcal{D}) = \theta^{\#[H]}(1 - \theta)^{\#[T]}$$

where  $\#[H]$  and  $\#[T]$  denote the number of heads and tails in the training data, respectively

# Maximum Likelihood Parameter Estimation

$$L(\boldsymbol{\theta} : \mathcal{D}) = \theta^{\#[H]}(1 - \theta)^{\#[T]}$$

- In practice, it is convenient to maximize the log-likelihood

$$\begin{aligned}\ell(\boldsymbol{\theta} : \mathcal{D}) &= \log \left( \theta^{\#[H]}(1 - \theta)^{\#[T]} \right) \\ &= \#[H] \log \theta + \#[T] \log(1 - \theta)\end{aligned}$$

- Taking the gradient and setting it to zero yields

$$\frac{\#[H]}{\theta} - \frac{\#[T]}{1 - \theta} \implies \theta = \frac{\#[H]}{\#[H] + \#[T]}$$

which is simply the fraction of heads

# Maximum Likelihood Parameter Estimation

- The *likelihood function* is the probability of the data

$$L(\boldsymbol{\theta} : \mathcal{D}) = \prod_m P(\mathbf{x}^{(m)}; \boldsymbol{\theta})$$

- A function  $\tau(x)$  is a *sufficient statistic* if, for any  $\mathcal{D}$  and  $\mathcal{D}'$  and  $\boldsymbol{\theta} \in \Theta$

$$\sum_m \tau(x^{(m)}) = \sum_{\bar{m}} \tau(x'^{(\bar{m})}) \Rightarrow L(\boldsymbol{\theta} : \mathcal{D}) = L(\boldsymbol{\theta} : \mathcal{D}')$$

- The tuple of counts  $\{\#[1], \dots, \#[K]\}$ , where  $\#[k]$  is the number of occurrences of  $x^k$  in  $\mathcal{D}$ , is a sufficient statistic for multinomials

$$L(\boldsymbol{\theta} : \mathcal{D}) = \prod_k \theta_k^{\#[k]}$$

- The maximum likelihood estimate for a multinomial is

$$\hat{\theta}_k = \frac{\#[k]}{M}$$



# ML Parameter Estimation: Sufficient Statistics

- Why do we refer to  $\tau(\mathbf{X})$  as a *sufficient statistic*?
- Sufficiency characterizes what is essential in a dataset
- A *statistic* is any function on the sample space that isn't a function of the parameter
- We say that  $\tau(\mathbf{X})$  is *sufficient* if there is no information in  $\mathbf{X}$  about  $\theta$  that isn't available in  $\tau(\mathbf{X})$
- Consider  $\theta$  to be a random variable
- In the Bayesian sense,  $\tau(\mathbf{X})$  is sufficient if

$$\theta \perp \mathbf{X} \mid \tau(\mathbf{X})$$

# ML Parameter Estimation: Sufficient Statistics

- Gaussian Distribution

$$P_{\boldsymbol{\theta}}(\mathbf{x}) = \frac{1}{Z(\boldsymbol{\theta})} \exp \left( \mathbf{t}(\boldsymbol{\theta})^\top \boldsymbol{\tau}(\mathbf{x}) \right) = \frac{1}{\sqrt{2\pi}\sigma} \exp \left( -\frac{(x - \mu)^2}{2\sigma^2} \right)$$

where

$$\boldsymbol{\tau}(x) = \begin{bmatrix} x & x^2 \end{bmatrix}$$

$$\boldsymbol{\theta} = \begin{bmatrix} \mu & \sigma^2 \end{bmatrix} \in \mathbb{R} \times \mathbb{R}^+$$

$$\mathbf{t}(\mu, \sigma^2) = \begin{bmatrix} \frac{\mu}{\sigma^2} & -\frac{1}{2\sigma^2} \end{bmatrix}$$

$$\implies \mu_{\text{MLE}} = \frac{1}{M} \sum_m \tau_1(x^{(m)}) = \frac{1}{M} \sum_m x^{(m)}$$

# ML Estimation for Bayesian Networks: Example

- Consider a Bayesian network  $X \rightarrow Y$ , where  $X$  and  $Y$  are binary
- The parameters defining the CPDs are

$$P(X) : \boldsymbol{\theta}_X = [\theta_{x^0}, \theta_{x^1}]$$

$$P(Y | X) : \boldsymbol{\theta}_{Y|x^0} = [\theta_{y^0|x^0}, \theta_{y^1|x^0}], \boldsymbol{\theta}_{Y|x^1} = [\theta_{y^0|x^1}, \theta_{y^1|x^1}]$$

- Assume we have a data set of samples  $\mathcal{D} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(M)}\}$
- The likelihood function follows as

$$\begin{aligned} L(\boldsymbol{\theta} : \mathcal{D}) &= \prod_m P(y^{(m)} | x^{(m)}; \boldsymbol{\theta}) P(x^{(m)}; \boldsymbol{\theta}) \\ &= \left( \prod_m P(y^{(m)} | x^{(m)}; \boldsymbol{\theta}_{Y|X}) \right) \left( \prod_m P(x^{(m)}; \boldsymbol{\theta}_X) \right) \end{aligned}$$

- Each term *decomposes* into separate *local likelihoods* for each variable, and is a function of the variable's CPD

# ML Estimation for Bayesian Networks: Example

$$L(\boldsymbol{\theta} : \mathcal{D}) = \left( \prod_m P(y^{(m)} | x^{(m)}; \boldsymbol{\theta}_{Y|X}) \right) \left( \prod_m P(x^{(m)}; \boldsymbol{\theta}_X) \right)$$

- The term  $\prod_m P(x^{(m)}; \boldsymbol{\theta}_X)$  is the likelihood for a multinomial
- The likelihood decomposes into a product of terms for each group of parameters

$$\prod_m P(y^{(m)} | x^{(m)}; \boldsymbol{\theta}_{Y|X}) = \prod_{m:x^{(m)}=x^0} P(y^{(m)} | x^{(m)}; \boldsymbol{\theta}_{Y|x^0}) \cdot \prod_{m:x^{(m)}=x^1} P(y^{(m)} | x^{(m)}; \boldsymbol{\theta}_{Y|x^1})$$

- As with the multinomial, we can write each as

$$\prod_{m:x^{(m)}=x^0} P(y^{(m)} | x^{(m)}; \boldsymbol{\theta}_{Y|x^0}) = \theta_{y^0|x^0}^{\#[x^0, y^0]} \cdot \theta_{y^1|x^0}^{\#[x^0, y^1]}$$

- The maximum likelihood parameters follow as

$$\theta_{y^0|x^0} = \frac{\#[x^0, y^0]}{\#[x^0]}$$

# ML Estimation for Bayesian Networks

- Suppose that we know the Bayesian network structure  $G$
- Let  $\theta_{X_i | \text{Pa}_{X_i}}$  be the parameters that determine the CPD  $P(X_i | \text{Pa}_{X_i})$
- Assume we have a data set of samples  $\mathcal{D} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(M)}\}$
- Maximum likelihood estimation corresponds to maximizing the log-likelihood  $\ell(\theta : \mathcal{D})$  (equivalent to maximizing the likelihood):

$$\begin{aligned} \frac{1}{M} \sum_{m=1}^M \log P(\mathbf{x}^{(m)}; \theta) &= \frac{1}{M} \sum_{m=1}^M \sum_{i=1}^N \log P(x_i^{(m)} | \text{Pa}_{X_i}; \theta_{X_i | \text{Pa}_{X_i}}) \\ &= \sum_{i=1}^N \frac{1}{M} \sum_{m=1}^M \log P(x_i^{(m)} | \text{Pa}_{X_i}; \theta_{X_i | \text{Pa}_{X_i}}) \end{aligned}$$

(subject to non-negativity and normalization constraints)

- **Global decomposability:** Likelihood decomposes into a product of independent terms, one for each set of parameters
- Results in an independent optimization problem for each CPD with a simple closed-form solution (objective is concave)


# ML Estimation for Bayesian Networks: Decomposability

$$\frac{1}{M} \sum_{m=1}^M \log P(\mathbf{x}^{(m)}; \boldsymbol{\theta}) = \sum_{i=1}^N \frac{1}{M} \sum_{m=1}^M \log P(x_i^{(m)} | \text{Pa}_{X_i}; \boldsymbol{\theta}_{X_i} | \text{Pa}_{X_i})$$

- Consider the four-node Bayesian network

$$P(\mathbf{X}; \boldsymbol{\theta}) = P(X_1; \theta_1) P(X_2 | X_1; \theta_2) P(X_3 | X_1; \theta_3) P(X_4 | X_2, X_3; \theta_4)$$

- Global decomposability allows us to break this up into four small Bayesian networks



# ML Estimation for Bayesian Networks

- Maximum likelihood estimation corresponds to maximizing:

$$\frac{1}{M} \sum_{m=1}^M \log P(\mathbf{x}^{(m)}; \boldsymbol{\theta}) = \sum_{i=1}^N \frac{1}{M} \sum_{m=1}^M \log P(x_i^{(m)} | \text{Pa}_{X_i}; \boldsymbol{\theta}_{X_i | \text{Pa}_{X_i}})$$

- Consider a tabular CPD and a random variable  $X$  with parents  $\mathbf{U}$

$$L_X(\boldsymbol{\theta}_{X|\mathbf{U}} : \mathcal{D}) = \prod_m \theta_{x^{(m)}|\mathbf{u}^{(m)}} = \prod_{\mathbf{u} \in \text{Val}(\mathbf{U})} \prod_{x \in \text{Val}(X)} \theta_{x|\mathbf{u}}^{\#[x, \mathbf{u}]}$$

- We can optimize each of the local likelihoods separately for each value of  $\mathbf{u}$ , which correspond to multinomial likelihoods

$$\hat{\theta}_{x|\mathbf{u}} = \frac{\#[x, \mathbf{u}]}{\#[\mathbf{u}]}$$

- The number of expected assignments for a particular parent decreases exponentially in the number of parents (known as *data fragmentation*)

# Limitations of ML Estimation

- Maximum likelihood estimation is purely data-driven
  - Consider estimating  $P(H)$  for a coin and a rubber duck
  - Suppose that we toss each 10 times and get 3 heads for each
  - MLE would assign  $P(H) = 0.3$  to both the coin and rubber duck
  - However, our intuition suggests that  $P(H) \approx 0.5$  for the coin
- MLE does not consider any prior knowledge that we might have about the parameters (i.e., in the form of a prior over the parameters)
- MLE does not provide a measure of confidence
  - Suppose that we had tossed the coin/duck 1000000 times and they had come up heads 300000 times each
  - Whether we use the experiment involving 10 tosses or the experiment involving 1000000 tosses, MLE estimates  $P(H) = 0.3$
  - MLE can't capture the fact that we would intuitively trust the larger experiment much more



# Bayesian Parameter Estimation

- A prior over the parameters  $\theta$  allows us to incorporate knowledge of the parameters (i.e., the parameters are now random variables)
- This gives rise to a joint model over the data and parameters

$$P(\mathcal{D}, \theta) = P(\mathcal{D} | \theta)P(\theta)$$


where  $P(\mathcal{D} | \theta)$  is the likelihood function from MLE and  $P(\theta)$  is the *parameter prior*

- Bayes' rule results in the *parameter posterior*

$$P(\theta | \mathcal{D}) = \frac{P(\mathcal{D} | \theta)P(\theta)}{P(\mathcal{D})}$$

where  $P(\mathcal{D}) = \int P(\mathcal{D} | \theta)P(\theta)d\theta$  is the *marginal likelihood*

# Bayesian Parameter Estimation: Example



- Let's return to the coin flip example, where the objective is to estimate  $\theta = P(H)$
- For MLE, where  $\theta$  is fixed, we assumed that the samples were marginally independent
- However, since  $\theta$  is a random variable, the samples are no longer marginally independent, since they provide information regarding  $\theta$
- Given  $\theta$ , the samples are conditionally independent

# Bayesian Parameter Estimation: Example

- Consider the joint distribution  $P(\mathcal{D}, \theta)$

$$\begin{aligned} P(x^{(1)}, \dots, x^{(M)}, \theta) &= P(x^{(1)}, \dots, x^{(M)} \mid \theta) P(\theta) \\ &= \left( \theta^{\# [H]} (1 - \theta)^{\# [T]} \right) P(\theta) \end{aligned}$$

- For Bayesian estimation, we are interested in the parameter posterior

$$P(\theta \mid x^{(1)}, \dots, x^{(M)}) \propto P(x^{(1)}, \dots, x^{(M)} \mid \theta) P(\theta)$$

which is a product of a Bernoulli distribution and the parameter prior

- We would like the prior and posterior to be of the same family, ideally with a closed-form update to the distribution's hyperparameters (we want a family that is the *conjugate prior* for the likelihood function)

# Bayesian Parameter Estimation: Example

- A **Beta distribution** is parameterized by two parameters  $\alpha_1, \alpha_0 \in \mathbb{R}^+$

$$\theta \sim \text{Beta}(\alpha_1, \alpha_0) \text{ with } p(\theta) = \gamma \theta^{\alpha_1-1} (1-\theta)^{\alpha_0-1}$$

where  $\gamma = \frac{\Gamma(\alpha_1+\alpha_0)}{\Gamma(\alpha_1)\Gamma(\alpha_0)}$  and  $\Gamma(x) = \int_0^\infty t^{x-1} e^{-t} dt$  is the *Gamma function*

- Consider the marginal distribution over the first coin toss

$$\begin{aligned} P(x^{(1)} = H) &= \int_0^1 P(x^{(1)} = H \mid \theta) P(\theta) d\theta \\ &= \int_0^1 \theta P(\theta) d\theta = \frac{\alpha_1}{\alpha_1 + \alpha_0} \end{aligned}$$

$\alpha_1$  and  $\alpha_0$  act like the number of imaginary Heads and Tails that we have thrown prior to sampling

# Bayesian Parameter Estimation: Example

- Let's consider the posterior after drawing  $M$  samples


$$\begin{aligned}P(\theta | x^{(1)}, \dots, x^{(M)}) &\propto P(x^{(1)}, \dots, x^{(M)} | \theta) P(\theta) \\&\propto \theta^{\#[H]} \cdot (1 - \theta)^{\#[T]} \cdot \theta^{\alpha_1 - 1} (1 - \theta)^{\alpha_0 - 1} \\&= \theta^{\#[H] - 1} (1 - \theta)^{\alpha_0 + \#[T] - 1} \\&= \theta^{\alpha_1 + \#[H] - 1} (1 - \theta)^{\alpha_0 + \#[T] - 1} \\&= \text{Beta}(\alpha_1 + \#[H], \alpha_0 + \#[T])\end{aligned}$$

- The Beta distribution is conjugate to the Bernoulli distribution
- The hyperparameter update is straightforward


$$\alpha_1 \leftarrow \alpha_1 + \#[H] \quad \alpha_0 \leftarrow \alpha_0 + \#[T]$$

i.e., we are essentially updating the counts for Heads and Tails

# Bayesian Parameter Estimation: Example



Beta(1, 1)



Beta(10, 10)

$$\alpha_1 \leftarrow \alpha_1 + \#[H] \quad \alpha_0 \leftarrow \alpha_0 + \#[T]$$

- Suppose that  $\mathcal{D}$  contains 3 Heads and 7 Tails
  - $P(\theta) = \text{Beta}(1, 1)$ :  $P(\theta | \mathcal{D}) = \text{Beta}(4, 8)$
  - $P(\theta) = \text{Beta}(10, 10)$ :  $P(\theta | \mathcal{D}) = \text{Beta}(13, 17)$
- Consider the likelihood of the next toss  $P(X^{(M+1)} = H | \mathcal{D}) = \frac{\alpha_1}{\alpha_1 + \alpha_0}$ 
  - $P(\theta) = \text{Beta}(1, 1)$ :  $P(X^{(M+1)} = H | \mathcal{D}) = 4/12 \approx 0.33$
  - $P(\theta) = \text{Beta}(10, 10)$ :  $P(X^{(M+1)} = H | \mathcal{D}) = 13/30 \approx 0.43$

# Bayesian Parameter Estimation: Example

- Recall the example of estimating  $P(H)$  for a coin and a rubber duck
- We tossed each 10 times and got 3 heads for each
- MLE assigned  $P(H) = 0.3$  to both the coin and rubber duck, which is inconsistent with intuition
- Instead we might use a Beta prior with  $\alpha_1 = \alpha_0 = 100$  for the coin and  $\alpha_1 = \alpha_0 = 1$  for the duck

# Bayesian Parameter Estimation: Prior Distribution

- We would like to express the posterior in terms of sufficient statistics just as we can the likelihood function, e.g., for a multinomial

$$L(\boldsymbol{\theta} : \mathcal{D}) = \prod_k \theta_k^{\# [k]}$$

- This depends upon the nature of the prior over parameters
- One effective prior is the *Dirichlet* distribution

$$\boldsymbol{\theta} \sim \text{Dirichlet}(\alpha_1, \dots, \alpha_K) \text{ s.t. } P(\boldsymbol{\theta}) \propto \prod_k \theta_k^{\alpha_k - 1}$$

where  $\{\alpha_1, \alpha_2, \dots, \alpha_K\}$  is a set of hyperparameters

- The Dirichlet is a generalization of the Beta distribution



# Bayesian Parameter Estimation: Conjugate Priors

- The Dirichlet is the conjugate prior to the multinomial, i.e., if the prior is Dirichlet (i.e.,  $P(\boldsymbol{\theta}) = \text{Dirichlet}(\alpha_1, \dots, \alpha_K)$ ), then the posterior is also Dirichlet,  $P(\boldsymbol{\theta} | \mathcal{D}) = \text{Dirichlet}(\alpha_1 + \#[1], \dots, \alpha_K + \#[K])$
- A conjugate prior makes it easy to update our distribution over the parameters, i.e., from

$$P(\boldsymbol{\theta}) = \text{Dirichlet}(\alpha_1, \dots, \alpha_K)$$


to

$$P(\boldsymbol{\theta} | \mathcal{D}) = \text{Dirichlet}(\alpha_1 + \#[1], \dots, \alpha_K + \#[K])$$

- Beyond the Dirichlet distribution, there are other conjugate priors for other distributions (e.g., Gaussians are self-conjugate)

# Bayesian Parameter Estimation for Bayesian Networks

- Consider estimating the parameters for a Bayesian network  $X \rightarrow Y$
- The associated parameters are  $\theta_X$  and  $\theta_{Y|X}$
- Since the parameters are now random variables, we can formulate the joint likelihood  $P(\mathcal{D}, \theta)$  as a Bayesian network



- Data instances  $(x^{(i)}, y^{(i)})$  and  $(x^{(j)}, y^{(j)})$  are independent given the latent parameters

# Bayesian Parameter Estimation for Bayesian Networks

- Consider again the parameter posterior

$$P(\boldsymbol{\theta} | \mathcal{D}) = \frac{P(\mathcal{D} | \boldsymbol{\theta})P(\boldsymbol{\theta})}{P(\mathcal{D})}$$

- The likelihood function decomposes into local likelihoods

$$P(\mathcal{D} | \boldsymbol{\theta}) = \prod L_i(\boldsymbol{\theta}_{X_i} | \text{Pa}_{X_i} : \mathcal{D})$$

- If *global parameter independence* holds (i.e., the prior decomposes),

$$P(\boldsymbol{\theta}) = \prod_i P(\boldsymbol{\theta}_{X_i} | \text{Pa}_{X_i})$$

- The posterior becomes

$$P(\boldsymbol{\theta} | \mathcal{D}) = \prod_i P(\boldsymbol{\theta}_{X_i} | \text{Pa}_{X_i} | \mathcal{D})$$